

## E.2 [GPT-2 XL only] Knowledge Neurons (KN)

The method by Dai et al. (2022) first selects neurons that are associated with knowledge expression via gradient-based attributions, and then modifies  $\text{mlp}_{proj}^{(l)}$  at the rows corresponding to those neurons by adding scaled embedding vectors. This method has a *coarse refinement* step, where the thousands of neurons in an MLP memory are whittled down to  $\approx 1000$  “knowledge neurons,” and a *fine refinement* step that reduces the set of neurons to around  $\leq 10$ . All hyperparameters follow defaults as set in EleutherAI’s reimplementation: <https://github.com/EleutherAI/knowledge-neurons>.

## E.3 [GPT-2 XL only] Knowledge Editor (KE)

De Cao et al. (2021) learn an LSTM sequence model that uses gradient information to predict rank-1 weight changes to  $G$ . Because the official code does not edit GPT-2, we use Mitchell et al. (2021)’s re-implementation in their study. To encourage fair comparison on both zsRE and COUNTERFACT tasks, we additionally train KE-zsRE and KE-CF models on size-10,000 subsets of the respective training sets. Hyperparameters for training are adopted from the given default configuration. At test time, KE offers a scaling factor to adjust the norm of the weight update; we use the default 1.0.

## E.4 [GPT-2 XL, GPT-J] Model Editor Networks with Gradient Decomposition (MEND)

Mitchell et al. (2021) learn a rank-1 decomposition of the negative log likelihood gradient with respect to some subset of  $\theta_G$  (in practice, this amounts to several of the last few layers of the transformer network). Again, for fair comparison, we train new versions of MEND (MEND-zsRE, MEND-CF) on the same sets that KE-zsRE and KE-CF were trained on. Similar to KE, hyperparameters for training and test-time inference are adopted from default configurations.

## E.5 [GPT-2 XL, GPT-J] Rank-One Model Editing (ROME)

ROME’s update (Section 3.1) consists of key selection (Eqn. 3), value optimization (Eqn. 4), and  $v$  insertion (Appendix A). We perform the intervention at layer 18. As Figure 1k shows, this is the center of causal effect in MLP layers, and as Figure 3 shows, layer 18 is approximately when MLP outputs begin to switch from acting as keys to values.

**Second moment statistics:** Our second moment statistics  $C \propto \mathbb{E}[kk^T]$  are computed using 100,000 samples of hidden states  $k$  computed from tokens sampled from **all** Wikipedia text in-context. Notice that sampling is not restricted to only special subject words; every token in the text is included in the statistic. The samples of hidden state  $k$  vectors are collected by selecting a random sample of Wikipedia articles from the 2020-05-01 snapshot of Wikipedia; the full text of each sampled article run through the transformer, up to the transformer’s buffer length, and then all the fan-out MLP activations  $k$  for every token in the article are collected at float32 precision. The process is repeated (sampling from further Wikipedia articles without replacement) until 100,000  $k$  vectors have been sampled. This sample of vectors is used to compute second moment statistics.

**Key Selection:** We sample 20 texts to compute the prefix ( $x_j$  in Eqn. 3): ten of length 5 and ten of length 10. The intention is to pick a  $k_*$  that accounts for the different contexts in which  $s$  could appear. Note that we also experimented with other  $x_j$  sampling methods:

- **No prefix.** This baseline option performed worse ( $S' = 86.1$  compared to  $S = 89.2$ ).
- **Longer prefixes.** Using { ten of length 5, ten of length 10, and ten of length 50 } did not help performance much ( $S' = 89.3$ ).
- **More same-length prefixes.** Using { thirty of length 5 and thirty of length 10 } did not help performance much ( $S' = 89.2$ ).

**Value Optimization:**  $v_*$  is solved for using Adam with a learning rate of 0.5 and  $1.5 \times 10^{-3}$  weight decay. The KL divergence scaling factor, denoted  $\lambda$  in Eqn. 4, is set to  $1 \times 10^2$ . The minimization loop is run for a maximum of 20 steps, with early stopping when  $\mathcal{L}(z)$  reaches  $5 \times 10^{-2}$ .

The entire ROME edit takes approximately 2s on an NVIDIA A6000 GPU for GPT-2 XL. Hypernets such as KE and MEND are much faster during inference (on the order of 100ms), but they require hours-to-days of additional training overhead.